

The Relational Model (Part I)

Annoucement

- **Next Tuesday (9/20) class is cancelled**
 - Instructor on conference travel

Steps in Database Design


- Requirements Analysis
 - user needs; what must database do?
- Conceptual Design
 - high level description: ER
- Logical Design
 - translate ER into DBMS data model
- Schema Refinement
 - consistency, normalization

Today's lecture


- **Introduction to DBMS data model**
 - Concepts (table, schema, row, column, ...)
 - Using SQL to create table, add and delete tuples
 - Integrity constraints

Six Types of Databases


Relational




Analytical (OLAP)




Key-Value




Column-Family



Graph




Document




Six Types of Databases

CS442
focus


Relational




Analytical (OLAP)




Key-Value




Column-Family



Graph



Document



Copyright Kelly-McCreary & Associates, LLC

11

Relational Database: Definition

- ***Relational database: a set of relations.***
- ***Relation: made up of 2 parts:***
 - *Schema* : specifies the name and *attributes* of relation
 - *Instance*
 - A *table* with rows and columns.
 - Consistent with schema

Relational Schema

- A *Schema* for a relation is represented in the format:
relation_name (attr1:type, ... attrn:type)
 - Example:
Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *GPA*: real)
- Attributes are referenced by name, not column locations
- Attribute names must be unique

Exercise



Student	Course	Course
Anna	DB Mgt.	Operating System
Bob	DB Mgt.	Web Programming
Cathy	Operating System	Artificial Intelligence

- **Is this table a valid relational table?**

Exercise (Cont.)



Student	Course1	Course 2
Anna	DB Mgt.	Operating System
Bob	DB Mgt.	Web Programming
Cathy	Operating System	Artificial Intelligence

- **Is this table a valid relational table?**

Relational Instances

- Instance: a *table*, with rows and columns
- Attributes (or fields) are stored in columns.
- Tuples (or records) are stored in rows.
- Attributes have a domain – an atomic type.
- The cardinality of the relation R = the number of rows in R (excludes the first row!)
- The degree/arity of the relation R = the number of columns in R

Notes of Relational Model

- No duplicate tuples in a relation
 - Question: what if we want to insert duplicate tuples? What can we do?
- Ordering
 - No ordering of tuples in a relation
- The value of each attribute is either drawn from its domain or the special value *NULL*
- Attribute's values are *atomic*.



Example Instance of Students Relation

Schema

Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *GPA*: real)

Instance

SID	Name	Login	Age	GPA
53666	Jones	Jones@cs	18	3.4
53668	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- Cardinality = 3, degree/arity = 5.

Questions:

- (1) Do the values in each column have to be distinct?
- (2) Does each tuple (i.e., each row) have to be distinct?

Today's lecture

- **Introduction to DBMS data model**
 - Concepts (table, schema, row, column, ...)
 - Using SQL to create table, add and delete tuples
 - Integrity constraints

Defining a Relation Schema in SQL

- **SQL (pronounced SEQUEL): standard language to describe and manipulate relational database**
 - Data Definition Language (DDL)
 - Create, modify, delete relations
 - Specify constraints
 - Administer users, security, etc.
 - Data Manipulation Language (DML)
 - Specify *queries* to find tuples that satisfy criteria
 - Add, modify, remove tuples

Creating Schema in SQL

- SQL syntax

```
CREATE TABLE table_name (  
    field1 datatype,  
    field2 datatype,  
    ...  
);
```


Data Types in CREATE TABLE Statement

All attributes must have a data type

1) **Character**: strings of fixed or varying length

CHAR(*n*): a fixed-length string of *n* character

VARCHAR(*n*): a set of characters of up to *n* characters.

2) **BOOLEAN**: an attribute whose value is logical

- The possible values are TRUE, FALSE, and UNKNOWN

3) **DATE & TIME**: represent dates and times

Data Types (Cont.)

- 4) **INT or INTEGER**: denotes typical integer values
- 5) **FLOAT**: denotes floating-point numbers
 - Real numbers with a fixed decimal point
 - **DECIMAL (n,d)** allows values that consists of *n* decimal digits, with the decimal point assumed to be *d* positions from the right

Example: Creating Schema in SQL

- **SQL syntax**

CREATE TABLE <name> (<field> <domain>, ...)

- **Example: creates the Students relation.**

Schema

Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *GPA*: real)

```
CREATE TABLE Students
    (sid CHAR(20),
     name CHAR(20),
     login CHAR(10),
     age INTEGER,
     GPA FLOAT);
```

Note: the type
(domain) of each field
is specified, and
enforced by the DBMS



Creating Schema in SQL - Exercise

Write the SQL statement that creates the *Enrolled* table.

Schema

Enrolled (*sid*: string, *cid*: string, *grade*:real)

Case-sensitivity of Table/Column Names

- **Is the table/column name in SQL statement case sensitive?**
 - By default,
 - Case-sensitive on Linux
 - Case-insensitive on Windows
 - MySQL has a configuration option to enable/disable it.

Creating Instances in SQL

SQL syntax

```
INSERT INTO table_name (field1, field1, ...)  
VALUES (value1, value2, ...);
```

Note: the order of columns and their values must match

- **Insert a single tuple**

```
INSERT INTO Students (SID, Name, Login, Age, GPA)  
VALUES ('53688', 'Smith', 'smith@ee', 18, 3.2);
```

- **Insert multiple tuples**

```
INSERT INTO Students (sid, name, login, age, GPA)  
VALUES ('53666', 'Jones', 'jones@cs', 18, 3.4),  
      ('53690', 'Smith', 'smith@eecs', 18, 3.2),  
      ('53650', 'Smith', 'smith@math', 19, 3.8);
```

Creating Instances in SQL (Cont.)

- Insert data only in specified columns

```
INSERT INTO table_name (field1, field1, ...)  
VALUES (value1, value2, ...)
```

Example: INSERT INTO Students (SID, Name)
VALUES('53688', 'Smith');

Note: key attributes must have inserted values.

- Insert data in all columns
 - Do not need to specify the column names in the statement.

```
INSERT INTO table_name  
VALUES (value1, value2, value3, ...);
```

Example: INSERT INTO Students
VALUES('53688', 'Smith', 'smith@ee', 18, 3.5);

- Make sure the order of the values is in the same order as the columns in the table.

Today's lecture

- **Introduction to DBMS data model**
 - Concepts (table, schema, row, column, ...)
 - Using SQL to create schema and table
 - Integrity constraints

Integrity Constraints

- **Integrity constraints (ICs): conditions specified on a database schema**
- **Types of ICs**
 - Keys
 - Foreign keys
 - Domain constraints: (e.g., the age of the driver license holders must be at least 16)

Superkey, Key, Primary Key, Candidate Key

- Similar to superkey, key, and primary key in ER diagram
- **Superkey:** A set of fields is a superkey if no two distinct tuples have same values in these key fields
 - *ALL* attributes of a relation together form a super key for the relation
- **Key:** Minimal superkey (no subset of the key is a superkey)
- **Primary key** and **candidate keys:**
 - >1 key for a relation?
 - One of the keys is chosen (by DBA) to be the primary key.
 - The other keys are called as candidate keys

Non-NULL Key Values in Relational Models

- All the keys (primary and candidate keys cannot contain a NULL value on any of their attributes
 - E.g., if the key is defined as a composite key (A, B), all records must have values on both A and B attributes.
 - Nulls value on either attribute A or B is not allowed.

Multiple Keys

- An instance can have multiple keys.
 - For example, the *student* table can have 3 keys: (SID), (Name, BirthDate), and (Name, DormRoomNumber)
- Composite keys can overlap
 - For example, (Name, BirthDate) and (Name, DormRoomNumber)
- Question: can the composite keys and singleton keys overlap?

Exercise

SID	Name	Login	Age	GPA
53666	Jones	Jones@cs	18	3.4
53668	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- **Question 1:** for the relation above, can the attribute *Name* be used as a key? What about *Login*? What about *(Name, Age)*?
- **Question 2:** for the relations below, are the following keys correct?

CID	SID	Grade
CS442	27001	A
CS510	27001	B
CS443	21097	A
CS510	24331	B

Key1: (CID)
Key2: (SID)
Key3: (CID, Grade)
Key4: (CID, SID, Grade)

Name	Grade
Alice.S	C
Alan.B	B
Adam.H	A
Alan.B	B

Key1: (Name)
Key2: (Name, Grade)

Nice Properties of Keys

Given a set of attributes A (can have 1 or multiple attributes):


- If A is a key, then any superset A' of A ($A \subseteq A'$) CANNOT be a key (*because key must be minimal*)
 - E.g., if (SSN) is a key, then any superset of SSN is not a key
- If A is not a key, then any subset A' of A ($A' \subseteq A$) CANNOT be a key.
 - E.g., if (Name, BirthDate) is not a key, then neither (Name) nor (BirthDate) is a key

How to Find *All* Keys Efficiently?

- Naïve way: enumerate all possible attribute
 - E.g., A table with 3 attributes (A, B, C)
 - All possible attribute sets (7 in total): (A), (B), (C), (A, B), (A, C), (B, C), (A, B, C)
 - There are $2^k - 1$ such sets in total for k attributes
- An efficient way
 - Construct an *attribute lattice* (explained in the next slide) to enumerate all possible combinations of attributes
 - Use the two properties of keys (previous slide) when traversing the attribute lattice to find keys.


Attribute Lattice

- Top of lattice: each node contains a single attribute
- Bottom of lattice: one node that contains all attributes
- The *i*th level of lattice: each node contains *i* attributes
- Edge $N1 \rightarrow N2$: $N1$ is a strict subset of $N2$



An example of attribute lattice of 3 attributes {A, B, C}

Finding *ALL* Keys from Attribute Lattice



- Starting from the top of lattice, for each node N
 - Check if N can be a key.
 - If N is a key, remove all descendants of N from lattice (because they cannot be keys as they are not minimal).
- Continue until all nodes in the lattice are traversed.



Exercise

- Consider the following facts, list ALL keys.
 - Each student has a student ID (SID), his/her name, and age;
 - Each student has a unique SID;
 - Some students can have the same name;
 - Some students can have the same age;
 - No two students have the same name and the same age.

Define Primary Keys by SQL

SQL Syntax for Primary Key definition:

```
CREATE TABLE table_name (  
    field1 datatype,  
    field2 datatype,  
    ...  
    PRIMARY KEY (att1, att2...)  
);
```

Notes:

- The key can contain either 1 attribute or multiple attributes
- Always put the key attributes in parentheses, even when there is only one attribute.

Example: Define Keys

- **Example: creates the Students relation.**

Schema

Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *GPA*: real)

```
CREATE TABLE Students
    (sid CHAR(20),
     name CHAR(20),
     login CHAR(10),
     age INTEGER,
     GPA FLOAT,
     PRIMARY KEY (sid));
```

Define Candidate Keys in SQL

```
CREATE TABLE table_name (  
    field1 datatype,  
    field2 datatype,  
    ...  
    PRIMARY KEY (k1_att1, k2_att2...),  
    UNIQUE (ck1_att1, ck1_att2...),  
    UNIQUE (ck2_att1, ck2_att2...)  
    ...  
);
```

Notes:

- K candidate keys should have k UNIQUE statements, one statement per candidate key.
- Always put the key attributes in parentheses, even when there is only one attribute.

A Simpler Way to Define Singleton Keys

Singleton keys (either primary or candidate keys)

- Place PRIMARY KEY or UNIQUE right after the type in the declaration of the key attribute.
- This is only allowed for singleton keys.
 - Composite keys only can be defined by using PRIMARY KEY and UNIQUE statements.
- Example:

```
CREATE TABLE Students (  
    SID CHAR(20) PRIMARY KEY,  
    SSN CHAR(9) UNIQUE,  
    age CHAR(20)  
);
```

Exercise

sid	cid	grade
10001	CS442	A
10002	CS442	B
10001	CS510	A



- Given schema: *Enrolled* (*sid*: string, *cid*: string, *grade*: real)
- Facts:
 - Each student has a unique sid.
 - Each course has a unique cid.
 - Each student can enroll in multiple courses. But each student can take the same course only once.
 - No two students in the same course receive the same grade.
- Questions:
 - What are the keys? (tips: there are 2 keys)
 - Write the SQL statement to define the *Enrolled* table with the key constraints.