


# SQL: The Query Language

## Part II

### R&G - Chapter 5



# SQL

- The form:

**SELECT  $A_1, A_2, \dots, A_n$   
FROM  $r_1, r_2, \dots, r_m$   
WHERE  $P$**

- $A_i$  represents an attribute
  - $r_i$  represents a relation
  - $P$  is a predicate
- This query is equivalent to the relational algebra expression:

$$\Pi_{A_1, A_2, \dots, A_n} (\sigma_P (r_1 \times r_2 \times \dots \times r_m))$$

# Roadmap of Today's lecture

- **Set operations**
  - Union
  - Intersect
  - Except

# Union

**Union:** R U S

In SQL:

Subquery 1

**UNION**

Subquery 2

- The two subqueries must be a valid SQL query (SELECT-FROM block, with WHERE and other clauses optional)
- UNION statement must be union-compatible:
  - The two subqueries must return the same attributes in SELECT clause

# UNION with Duplicate Rows

- **UNION excludes duplicate rows**

A
a1
a2
a2
a3

T1

A
a2
a3
a3

T2

```
SELECT A  
FROM T1  
UNION  
SELECT A  
FROM T2
```

The query returns

A
a1
a2
a3

# UNION of Two SELECT \* Clauses

A	B
a1	b1
a2	b1
a3	b2

T1

A	B
a1	b2
a2	b2

T2

```
SELECT *  
FROM T1  
UNION  
SELECT *  
FROM T2
```

Both T1 and T2 have A as the key

The query returns

A	B
a1	b1
a1	b2
a2	b1
a2	b2
a3	b2

The records of the same key but different non-key values are considered as different and add into the union result

# Example of Union

- **Schema**

- Boats (bid, bname, color)
- Sailors(sid, sname, rating, age)
- Reserves(sid, bid, day)

**Query 1: Find ID of sailors who've reserved a red or a green boat**

Solution 1 (without set operations)

$$\pi_{sid}(\sigma_{color='red' \vee color='green'} Boats \bowtie Reserves)$$

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND (B.color='red' OR B.color='green');
```

**Note:** AND is always processed before OR. Without () around `B.color='red' OR B.color='green'`, the WHERE clause will be evaluated as:  
`(R.bid=B.bid AND B.color='red') OR B.color='green';`

# Example of Union

- **Schema**

- Boats (bid, bname, color)
- Sailors(sid, sname, rating, age)
- Reserves(sid, bid, day)

**Query 1: Find ID of sailors who've reserved a red or a green boat**

**Solution 2 (with set union operation)**

```
SELECT  R.sid
FROM    Boats B NATURAL JOIN Reserves R
WHERE   B.color='red'
UNION
SELECT  R.sid
FROM    Boats B NATURAL JOIN Reserves R
WHERE   B.color='green' ;
```



# INTERSECT

**Intersection:**  $R \cap S$

In SQL:

Subquery 1

**INTERSECT**

Subquery 2

- The two subqueries must be a valid SQL query (SELECT-FROM block, with WHERE and other clauses optional)
- INTERSECT must be union-compatible:
  - The two subqueries must have the same attributes in SELECT clause

# INTERSECT with Duplicate Rows

- **INTERSECT** excludes duplicate rows

A
a1
a1
a2
a3

T1

A
a2
a3
a3

T2

```
SELECT A
FROM T1
INTERSECT
SELECT A
FROM T2
```

The query returns

A
a2
a3

# Example of Intersection

- **Schema**

- Boats (bid, bname, color)
- Sailors(sid, sname, rating, age)
- Reserves(sid, bid, day)

**Query 2: Find ID of sailors who've reserved a red and a green boat**  
**Solution 1 (with set operation)**

```
SELECT sid
FROM Boats B NATURAL JOIN Reserves R
WHERE B.color='red'
INTERSECT
SELECT sid
FROM Boats B NATURAL JOIN Reserves R
WHERE B.color='green' ;
```

## Query 2: Find ID of sailors who've reserved a red and a green boat (continue)

- Can we write equivalent SQL statement WITHOUT using set operations?
- Is the following solution correct?

```
SELECT R.sid  
FROM Boats B NATURAL JOIN Reserves R  
WHERE B.color='red' AND B.color='green'
```

Query 2: Find ID of sailors who've reserved a red and a green boat

Is the following solution correct?

```
SELECT R.sid  
FROM Boats B NATURAL JOIN Reserves R  
WHERE B.color='red' AND B.color='green'
```

NO single boat of both red and green colors. The query will return empty answer.

## Query 2: Find ID of sailors who've reserved a red and a green boat



- **Schema**

- Boats (bid, bname, color)
  - Sailors(sid, sname, rating, age)
  - Reserves(sid, bid, day)
- 
- Can we write equivalent SQL statement without using set operations?
    - Hint: join with Boats table TWICE (one for red boat, one for green boat)

## Query 2: Find ID of sailors who've reserved a red and a green boat



- **Schema**

- Boats (bid, bname, color)
  - Sailors(sid, sname, rating, age)
  - Reserves(sid, bid, day)
- Can we write equivalent SQL statement without using set operations?
    - Hint: join with Boats table TWICE (one for red boat, one for green boat)

```
SELECT R1.sid
FROM Boats B1, Boats B2, Reserves R1, Reserves R2
WHERE   B1.color='red'    //B1 only has red boats
        AND B2.color='green'; //B2 only has green boats
        AND R1.bid=B1.bid    // Natural join R1&B1 (resv. of red boats)
        AND R2.bid=B2.bid    // Natural join R2&B2 (resv. of green boats)
        AND R1.sid=R2.sid    // Reservation of same sailor
```

## Query 2: Find ID of sailors who've reserved a red and a green boat



- **Schema**

- Boats (bid, bname, color)
  - Sailors(sid, sname, rating, age)
  - Reserves(sid, bid, day)
- Can we write equivalent SQL statement without using set operations?
    - **Question:** can we use one single Reserves table in the two joins (as shown below)?

```
SELECT R.sid
FROM Boats B1, Boats B2, Reserves R
WHERE      R.bid=B1.bid
          AND R.bid=B2.bid
          AND B1.color='red'
          AND B2.color='green' ;
```



## Query 3: Find name of sailors who've reserved at least 2 different boats

- **Schema**
  - Boats (bid, bname, color)
  - Sailors(sid, sname, rating, age)
  - Reserves(sid, bid, day)
- **Write the query without using aggregate function**
- Hint: join with Boats table TWICE (for two different reservations)

```
SELECT S1.sname
FROM Reserves R1, Reserves R2, Sailors S1, Sailors S2
WHERE R1.sid=R2.sid  // same sailor
      AND R1.bid<>R2.bid  //Two different boats
      AND R1.sid=S1.sid  //natural join R1 & S1
      AND R2.sid=S2.sid; //natural join R2 & S2
```

# EXCEPT

**Set difference:  $R - S$**

*In SQL:*

Subquery 1

**EXCEPT**

Subquery 2

- The two subqueries must be a valid SQL query (SELECT-FROM block, with WHERE and other clauses optional)
- EXCEPT must be union-compatible:
  - The two subqueries must have the same attributes in SELECT clause

# EXCEPT

- EXCEPT (sometimes also use MINUS)
- EXCEPT only takes the distinct rows of queries

A
a1
a1
a2
a2
a3

T1

A
a2
a3
a3

T2

```
SELECT A  
FROM T1  
EXCEPT  
SELECT A  
FROM T2
```

The query returns

A
a1

# Example of EXCEPT

Query 4: find ID of sailors who've reserved a red boat  
**but never reserved** a green boat

```
SELECT sid
FROM Boats B NATURAL JOIN Reserves R
WHERE B.color='red'
EXCEPT
SELECT sid
FROM Boats B NATURAL JOIN Reserves R
WHERE B.color='green'
```

## Query 5: Find name of sailors who never reserved a boat

- **Schema**

- Boats (bid, bname, color)
- Sailors(sid, sname, rating, age)
- Reserves(sid, bid, day)

Is this solution correct?

```
SELECT sname
FROM Sailors
EXCEPT
SELECT sname
FROM Sailors NATURAL JOIN Reserves;
```

## Query 5: Find name of sailors who never reserved a boat

- **Schema**

- Boats (bid, bname, color)
- Sailors(sid, sname, rating, age)
- Reserves(sid, bid, day)

Is this solution correct?

```
SELECT sname
FROM Sailors
EXCEPT
SELECT sname
FROM Sailors NATURAL JOIN Reserves
```

- The solution is incorrect, as the sailors may have duplicated names, while some reserved boats, some don't.
- How to fix?

## Query 5: find name of sailors who never reserved a boat

```
CREATE TABLE Temp_Sid AS
SELECT sid
FROM Sailors
EXCEPT
SELECT sid
FROM Sailors NATURAL JOIN Reserves;

SELECT sname
FROM Temp_Sid NATURAL JOIN Sailors;
```

Learned lesson:

- Be careful when applying projection on non-key attributes. Duplicates may be trouble makers!